



HF



# Hugging Face

## The Comprehensive Guide

*From Transformers to Pipelines: A Visual Journey*

**Author: Mausam Kar**

Written for Beginners & Practitioners

# Contents

<b>1</b>	<b>Getting Started with Hugging Face</b>	<b>1</b>
1.1	What is Hugging Face? . . . . .	1
1.2	The Pipeline Abstraction . . . . .	1
1.2.1	How it Works . . . . .	1
1.3	Your First Code Example . . . . .	2
1.4	Summary . . . . .	2
<b>2</b>	<b>Models and Tokenizers</b>	<b>3</b>
2.1	Introduction . . . . .	3
2.2	The AutoClasses . . . . .	3
2.3	Tokenizers: Speaking the Language of Machines . . . . .	3
2.3.1	The Tokenization Pipeline . . . . .	3
2.4	Using AutoTokenizer . . . . .	4
2.5	Models and Configurations . . . . .	4
2.6	Saving Your Work . . . . .	5
2.7	Summary . . . . .	5
<b>3</b>	<b>The Datasets Library</b>	<b>6</b>
3.1	Introduction . . . . .	6
3.2	Loading Data . . . . .	6
3.2.1	From the Hub . . . . .	6
3.2.2	From Local Files . . . . .	6
3.3	Inspecting Data with Pandas . . . . .	6
3.4	Processing Data: Slicing and Mapping . . . . .	7
3.4.1	Slicing and Shuffling . . . . .	7
3.4.2	The Map Function . . . . .	7
3.5	Streaming: Handling Big Data . . . . .	7
3.6	Summary . . . . .	8
<b>4</b>	<b>Fine-Tuning Your First Model</b>	<b>9</b>
4.1	The Concept of Transfer Learning . . . . .	9
4.2	The Trainer API . . . . .	9
4.2.1	1. Computing Metrics . . . . .	9
4.2.2	2. Training Arguments . . . . .	10
4.2.3	3. Launching Training . . . . .	10
4.3	Checkpoints and Resuming . . . . .	10
4.4	Summary . . . . .	10

Contents	2
----------	---

---

<b>5 The Hugging Face Hub</b>	<b>11</b>
5.1 The GitHub of Machine Learning . . . . .	11
5.2 Authentication: The CLI . . . . .	11
5.3 Model Cards: Documentation Matters . . . . .	11
5.3.1 Metadata tags . . . . .	12
5.4 Hugging Face Spaces . . . . .	12
5.4.1 SDKs: Gradio vs Streamlit . . . . .	12
5.5 Conclusion . . . . .	12

# Chapter 1

## Getting Started with Hugging Face

### 1.1 What is Hugging Face?

Hugging Face is often referred to as the “**GitHub of Machine Learning**”. It is an open-source platform that provides tools to build, train, and deploy machine learning models, primarily centered around *Natural Language Processing (NLP)*.

At its core, Hugging Face democratizes AI by making state-of-the-art models accessible to everyone through the `transformers` library.

#### Key Definition: Transformers

**Transformers** are a type of deep learning model designed to process sequential data, such as text. Introduced by Google in 2017 (in the paper *"Attention is All You Need"*), they have revolutionized NLP by allowing models to understand context better than previous architectures like RNNs.

### 1.2 The Pipeline Abstraction

One of the easiest ways to use Hugging Face is through the `pipeline()` function. It abstracts away the complex code required to process text and generate predictions.

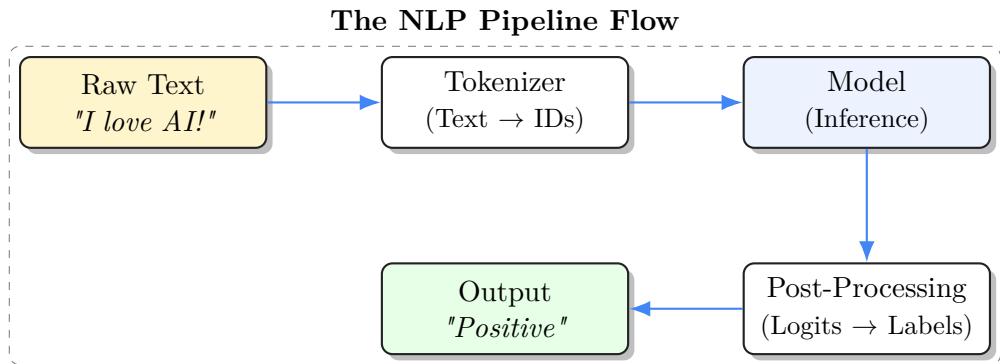
#### Concept

Think of a **Pipeline** as a factory assembly line. You pour raw material (text) into one end, and the finished product (sentiment, translation, summary) comes out the other end. The pipeline handles all the machinery inside.

#### 1.2.1 How it Works

The pipeline consists of three main steps:

1. **Preprocessing:** Converting text into numbers (Tokenization).
2. **Inference:** Running the numbers through the Model.
3. **Post-processing:** Converting the model’s output back into human-readable text.



### 1.3 Your First Code Example

Let's see how simple it is to use Python to analyze the sentiment of a sentence.

```
1 from transformers import pipeline
2
3 # 1. Initialize the pipeline
4 classifier = pipeline("sentiment-analysis")
5
6 # 2. Pass text to the pipeline
7 result = classifier("Hugging Face makes AI easy!")
8
9 # 3. Print the result
10 print(result)
11 # Output: [{'label': 'POSITIVE', 'score': 0.99}]
```

Listing 1.1: Basic Sentiment Analysis

#### Note

By default, the pipeline downloads a default model (like `distilbert`). You can specify other models if you need specific capabilities, such as multi-lingual support.

### 1.4 Summary

In this chapter, we learned:

- **Hugging Face** is a platform for sharing and using ML models.
- **Pipelines** simplify the process of using these models.
- We can write a sentiment analysis script in just 3 lines of code!

# Chapter 2

## Models and Tokenizers

### 2.1 Introduction

While the `pipeline` API is powerful, it is often a "black box." To truly master Hugging Face, you must understand the two fundamental components that power every NLP application: **Tokenizers** and **Models**.

This chapter delves into the `AutoClasses`, the process of tokenization, and how to manage model configurations and weights.

### 2.2 The AutoClasses

Hugging Face provides a set of classes designed to automatically select the correct architecture for a given checkpoint. This design pattern makes your code incredibly flexible and portable.

#### Key Classes

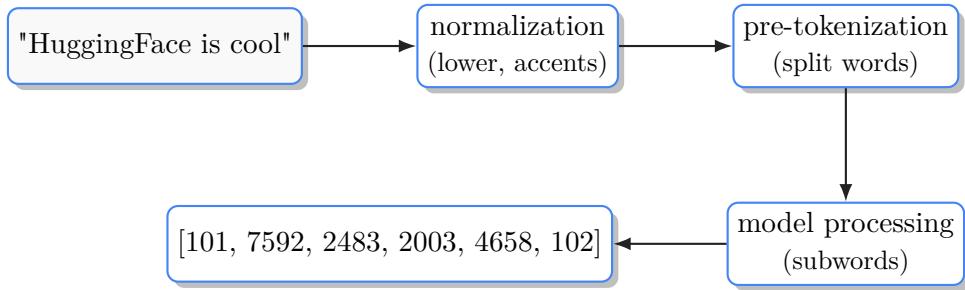
- `AutoTokenizer`: Handles text preprocessing.
- `AutoModel`: Loads the base model (transformer body).
- `AutoConfig`: Manages technical parameters (layers, hidden size).

### 2.3 Tokenizers: Speaking the Language of Machines

Deep Learning models cannot process raw strings. They require numerical input. The tokenizer bridges this gap by breaking text into smaller units called *tokens* and mapping them to integers (*Input IDs*).

#### 2.3.1 The Tokenization Pipeline

Tokenization isn't just splitting by spaces. It involves a sophisticated pipeline:



### Concept

**Subword Tokenization:** Algorithms like BPE (Byte-Pair Encoding) or WordPiece solve the "unknown token" problem. They split rare words into common subwords. Example: "unfriendly" → "un", "friend", "ly"

## 2.4 Using AutoTokenizer

Here is a complete example of loading a tokenizer and processing a batch of sentences.

```

1  from transformers import AutoTokenizer
2
3  # 1. Load a pre-trained tokenizer
4  checkpoint = "bert-base-cased"
5  tokenizer = AutoTokenizer.from_pretrained(checkpoint)
6
7  # 2. Define a batch of sentences
8  batch = [
9      "Hello, world!",
10     "Hugging Face course is amazing."
11 ]
12
13 # 3. Tokenize
14 # padding: pad short sequences to the longest in batch
15 # truncation: cut sequences longer than model max length
16 inputs = tokenizer(
17     batch,
18     padding=True,
19     truncation=True,
20     return_tensors="pt" # Return PyTorch tensors
21 )
22
23 print(inputs)
24 # Outputs dictionary with 'input_ids', 'attention_mask'
  
```

Listing 2.1: Processing a Batch

## 2.5 Models and Configurations

The `AutoModel` class loads the weights of the network. However, sometimes you want to inspect or modify the architecture before loading the weights. This is where `AutoConfig` comes in.

```
1 from transformers import AutoConfig, AutoModel
2
3 # Load the default configuration
4 config = AutoConfig.from_pretrained("bert-base-uncased")
5
6 # Inspect a parameter
7 print(config.hidden_size) # 768
8
9 # Modify it (e.g., for a smaller custom model)
10 config.num_hidden_layers = 10
11
12 # Initialize a model with this config (Random weights!)
13 model = AutoModel.from_config(config)
```

Listing 2.2: Customizing Configuration

### Note

Initializing from config loads **random weights**. You must use `from_pretrained()` to load trained knowledge.

## 2.6 Saving Your Work

After fine-tuning or modifying a model, you need to save it. Hugging Face makes this easy with `save_pretrained`.

```
1 # Save tokenizer and model to a local directory
2 save_directory = "./my_saved_model"
3
4 tokenizer.save_pretrained(save_directory)
5 model.save_pretrained(save_directory)
6
7 # You can now load from this directory!
8 loaded_model = AutoModel.from_pretrained(save_directory)
```

## 2.7 Summary

In this chapter, we explored:

- How **Tokenizers** transform text into input IDs using algorithms like BPE.
- The role of **AutoConfig** in defining model architecture.
- How to **Save and Load** models locally, ensuring your work is never lost.

# Chapter 3

## The Datasets Library

### 3.1 Introduction

Good models perform poorly on bad data. The Hugging Face `datasets` library is the standard tool for loading, processing, and sharing datasets in the NLP ecosystem. It is designed to be efficient, developer-friendly, and compatible with NumPy, Pandas, and PyTorch/TensorFlow.

### 3.2 Loading Data

Hugging Face hosts thousands of datasets on the Hub, but you will often work with your own local files.

#### 3.2.1 From the Hub

```
1 from datasets import load_dataset
2
3 # Load a benchmark dataset (GLUE)
4 glue_dataset = load_dataset("glue", "mrpc", split="train")
```

#### 3.2.2 From Local Files

The library supports loading directly from CSV, JSON, and text files.

```
1 # Load from a single CSV
2 data_files = {"train": "path/to/train.csv", "test": "path/to/test.
3               csv"}
4 dataset = load_dataset("csv", data_files=data_files)
5
6 # Load from JSON Lines
7 json_dataset = load_dataset("json", data_files="my_data.jsonl")
```

Listing 3.1: Loading Local Files

### 3.3 Inspecting Data with Pandas

Sometimes you just want to "see" your data. The `datasets` library integrates seamlessly with Pandas.

```

1 import pandas as pd
2
3 # Convert to pandas
4 df = glue_dataset.to_pandas()
5
6 # Show the first 5 rows
7 print(df.head())

```

Listing 3.2: Converting to Pandas DataFrame

**Efficiency Note**

The conversion to Pandas is purely in-memory. For massive datasets, you should avoid this or converting only a small slice.

## 3.4 Processing Data: Slicing and Mapping

### 3.4.1 Slicing and Shuffling

You can manipulate datasets just like python lists context.

```

1 # Shuffle the dataset
2 shuffled = dataset.shuffle(seed=42)
3
4 # Select the first 1000 examples
5 small_dataset = dataset.select(range(1000))

```

### 3.4.2 The Map Function

The core powerhouse of the library is `.map()`. It allows you to apply processing functions (like tokenization) to every element.

**Concept****Why `batched=True`?**

Tokenizers are written in Rust and can process lists of texts much faster than single strings thanks to parallelism. Enabling batching can speed up your processing by 10x or 100x.

```

1 def tokenize_function(examples):
2     # 'examples["text"]' is a LIST of strings
3     return tokenizer(examples["text"], truncation=True)
4
5 # Apply to the whole dataset
6 tokenized_dataset = dataset.map(tokenize_function, batched=True)

```

Listing 3.3: High-Performance Tokenization

## 3.5 Streaming: Handling Big Data

What if your dataset is 1TB and your laptop has 16GB of RAM? Enter **Streaming**.

Streaming allows you to iterate over a dataset without downloading it to disk. It streams data over the network on-the-fly.

```
1 # This returns an IterableDataset
2 c4 = load_dataset("c4", "en", streaming=True)
3
4 # Get the first example
5 print(next(iter(c4)))
```

## 3.6 Summary

We have covered:

- Loading data from the Hub and local CSV/JSON files.
- Interoperating with Pandas for visualization.
- Using `map()` for efficient preprocessing.
- Handling massive datasets with Streaming.

# Chapter 4

## Fine-Tuning Your First Model

### 4.1 The Concept of Transfer Learning

Traditional machine learning required training models from scratch on specific datasets. In the era of large language models, this is inefficient.

**Transfer Learning** allows us to take a model trained on a massive generic corpus (like Wikipedia) and "fine-tune" it on a small, task-specific dataset (like your company's emails).

#### The Head Analogy

Imagine the model as a body and a head.

- **Body:** Contains general language understanding (grammar, context). This is Frozen or fine-tuned slightly.
- **Head:** The final layer responsible for the output (e.g., "Positive/Negative"). This is **randomly initialized** and trained from scratch.

### 4.2 The Trainer API

The `Trainer` class is a high-level API that abstracts away the complex training loop of PyTorch (optimizers, backward pass, gradient accumulation).

#### 4.2.1 1. Computing Metrics

By default, the Trainer only logs the loss. To know if your model is actually working (Accuracy, F1), we need a `compute_metrics` function.

```
1 import numpy as np
2 import evaluate
3
4 metric = evaluate.load("accuracy")
5
6 def compute_metrics(eval_pred):
7     logits, labels = eval_pred
8     predictions = np.argmax(logits, axis=-1)
9     return metric.compute(predictions=predictions, references=
10                           labels)
```

### 4.2.2 2. Training Arguments

These control "how" the model learns.

```
1 from transformers import TrainingArguments
2
3 args = TrainingArguments(
4     output_dir='./results',
5     # Evaluation strategy: evaluate every epoch
6     evaluation_strategy="epoch",
7     # Learning rate: usually very small for Transfer Learning
8     learning_rate=2e-5,
9     per_device_train_batch_size=16,
10    num_train_epochs=3,
11    # Weight decay: regularization technique
12    weight_decay=0.01,
13 )
```

Listing 4.1: Essential Hyperparameters

### 4.2.3 3. Launching Training

Combine the model, args, data, and metrics into the Trainer.

```
1 from transformers import Trainer
2
3 trainer = Trainer(
4     model=model,
5     args=args,
6     train_dataset=tokenized_datasets["train"],
7     eval_dataset=tokenized_datasets["test"],
8     tokenizer=tokenizer,
9     compute_metrics=compute_metrics,
10 )
11
12 trainer.train()
```

## 4.3 Checkpoints and Resuming

Training can take hours. What if your computer crashes? The Trainer defines `save_steps` to save checkpoints automatically.

To resume training from the latest checkpoint:

```
1 trainer.train(resume_from_checkpoint=True)
```

## 4.4 Summary

Fine-tuning adapts a powerful general model to your specific needs. The `Trainer` API simplifies the training loop, evaluation, and checkpoint management, letting you focus on the data and the results.

# Chapter 5

## The Hugging Face Hub

### 5.1 The GitHub of Machine Learning

The Hugging Face Hub constitutes the central repository of the NLP community. It hosts over 300,000 models, 50,000 datasets, and thousands of demo applications (Spaces).

### 5.2 Authentication: The CLI

To interact with the Hub (uploading models), you need to authenticate. The command-line interface (CLI) is the easiest way to do this.

```
1 $ huggingface-cli login  
2  
3 # You will be asked to enter your token from:  
4 # https://huggingface.co/settings/tokens
```

Listing 5.1: Login via Terminal

#### Git Integration

Under the hood, every model on the Hub is a Git repository. You can `git clone`, `git add`, and `git push` large model files (using Git LFS) just like you would with code.

### 5.3 Model Cards: Documentation Matters

A model without documentation is useless. The "Model Card" is the `README.md` file of your repository. It contains:

- **Model Description:** What does it do?
- **Intended Use:** What are the valid use cases?
- **Limitations:** Bias and risks.
- **Training Data:** What data was it trained on?

### 5.3.1 Metadata tags

Model cards start with a YAML header that helps the Hub filter your model.

```
---
language: en
tags:
- text-classification
- pytorch
datasets:
- glue
metrics:
- accuracy
---
```

## 5.4 Hugging Face Spaces

Spaces are slightly different from Model repos. They are designed to host *running applications*.

### 5.4.1 SDKs: Gradio vs Streamlit

- **Gradio:** Built by Hugging Face. Extremely simple, "Python-function-to-UI" approach. Best for quick demos.
- **Streamlit:** More flexible, dashboard-oriented. Great for data exploration apps.

```
1 import gradio as gr
2 from transformers import pipeline
3
4 classifier = pipeline("sentiment-analysis")
5
6 def predict(text):
7     return classifier(text)[0]
8
9 # Launch the interface
10 gr.Interface(
11     fn=predict,
12     inputs="text",
13     outputs="label",
14     title="My Sentiment Analyzer"
15 ).launch()
```

Listing 5.2: A Full Gradio App

## 5.5 Conclusion

You have now journeyed from the basics of tokenization to training your own models and sharing them with the world. The Hugging Face ecosystem is vast, but you have the keys to explore it.

*"The magic of AI is not in the models, but in what you build with them."*